

Performance diagnostics in JS

What's ahead?

- How Node.js works
- Why performance goes bad
- Diagnosing!
- Perf improving principles

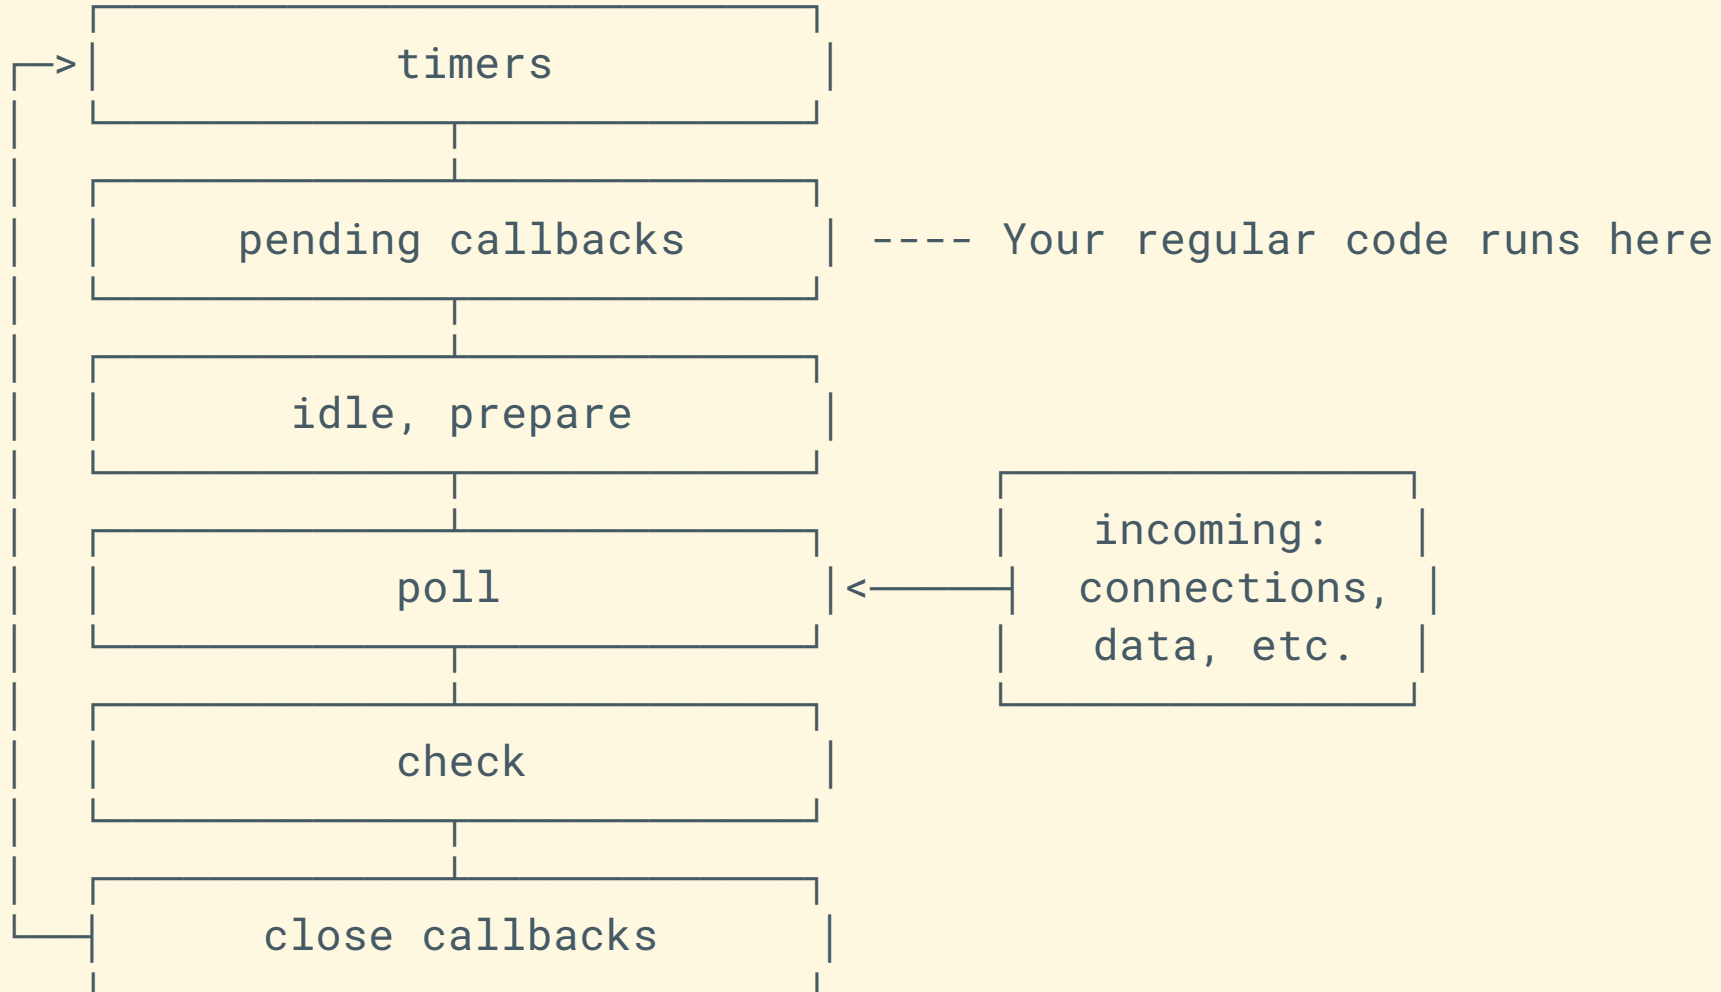
What's special about Node.js ?


It doesn't spawn threads for incoming requests

Consequences:

- Scales well on one process
- Offloads I/O to a pool of threads
 - Yes, there's threads, all 5 of them or so
- Asynchronous event loop model
 - Requests share memory and scope

Event loop





Classic event loop talk

<https://www.youtube.com/watch?v=8aGhZQkoFbQ>

<http://latentflip.com/loupe>

Performance, or what it isn't

Event loop blocked

Synchronous work going on means no other event loop functionalities work.

Incoming connections will wait on the network layer to get handled.

Heavy work examples

- JSON parsing
- String processing/concatenation
- Cryptography/math
- Checking if file exists (if failed) 🤔

Do not thrash the Node js event loop

<https://www.youtube.com/watch?v=VI29mUA8n9w>

Excess promises/async

- queues get filled up
- async steps, wrapping/unwrapping promises takes time

Memory issues and Garbage Collection

GC can take a lot of time if it has a lot to do!

- allocating lots of heavy short-lived objects puts pressure on GC

Programmer errors

Mostly just unnecessary iteration or duplicating work

- in 2014 or 2015 folks at Netflix discovered that updating express endpoints in a working app without restarting actually adds them to the routes array which makes routing slower and effectively leaks memory.

Clogging the thread pool

Mongodb injection

Collection.find(req.query)

```
?id=1  
?$where=sleep(20000)
```

Source: Node.js Interactive talk from my pal Vlad

<https://www.youtube.com/watch?v=xJWZsoYmsIE&list=PLfMzBWSH11xaZvhv1X5Fq1H-oMdnAtG6k&t=1104s>

Diagnosing

Finding event loop blocks

Detection:

`blocked` or `loopbench`

Measuring event loop utilization

```
require('perf_hooks').performance.eventLoopUtilization
```

<https://nodesource.com/blog/event-loop-utilization-nodejs/>

Finding the actual blocking code path:

`blocked-at` - it's an experiment of mine that needs refreshing 😅

CPU usage measurement

flame graphs

`0x` `clinic flame`

or Linux builtin `perf` with <https://nodejs.org/en/guides/diagnostics-flamegraph/> (might be outdated since I wrote it)

devtools

- Performance tab (recently simplified)

Other performance measurements

- trace opt/deopt
- V8 tracing overall
- the entire `perf_hooks`
- async hooks and promise hooks
- diagnostics channel

diagnostics_channel

https://www.youtube.com/watch?v=wqGlbxIHAwI&list=PL0CdG_OSSGIBYI7_e6Zs4kFSXL9LvOn8gM&index=4

Performance improving principles

Caching and memoization

- You can keep results in memory instead of calculating them. Works wonders for inefficient recursive graph traversal implementations as well as simple request-response sequences

Memoization is when you don't invalidate any caches, they get garbage collected at the end of the function execution or are expected to be finite

Caching is when you manage the storage.

Cache invalidation

Famously a hard problem. Just make sure you think twice.

Avoiding duplicate fetches

Wrong

```
const cache = new Map();

function getUrl(url) {
  if (cache.has(url)){
    return cache.get(url)
  } else {
    return fetch(url).then(r=>r.json())
      .then(result => {
        cache.set(url, result);
        return result;
      });
  }
}
```


Right (disregard the memory leak)

```
const cache = new Map();

function getUrl(url) {
  if (cache.has(url)){
    return cache.get(url)
  } else {
    const promiseToCache = fetch(url).then(r=>r.json());
    cache.set(url, promiseToCache);
    return promiseToCache;
  }
}
```

Avoiding duplicate fetches - tools

Use actual tools that have been benchmarked.

<https://github.com/mcollina/async-cache-dedupe>

Optimization and Hidden Classes

- monomorphic
- polymorphic
- megamorphic
- deopt is when compiled JS needs to be thrown away because types don't match anymore.

Avoid changing types of references and structures of objects

Fail fast

- return 503 so that load balancer redirects traffic if you see you're overwhelmed

Unexpected performance monsters

- config files lookup (especially if there's a tree traversal and checking for files all the time) - .babelrc for example
- some popular logging libraries are horribly slow by default

Other material

https://www.youtube.com/watch?v=wqGlbxIHAwI&list=PL0CdG_OSSGIBYI7_e6Zs4kFSXL9LvOn8gM&index=4

https://www.youtube.com/watch?v=J7q6HD06Sw0&list=PL0CdG_OSSGIBaULAdbribJiENfXxPW0aLQ&index=20

graphql specific stuff

2020 GraphQL - Accelerated! - Matteo Collina

<https://www.youtube.com/watch?v=CQ1KgfVcL2E>

2018 GraphQL - Accelerated! - Matteo Collina & Mathias Buus

<https://www.youtube.com/watch?v=-V86BB9wrG8>

graphql talk and workshop from NodeConfEU

Intro talk:

<https://www.youtube.com/watch?v=guAMBFqKSF4>

Workshop:

<https://github.com/nearform/the-graphql-workshop>

memory leaks etc

<https://naughtur.pl/pres3/mem2020/2024.html>